



UPPSALA UNIVERSITET

Computer-Assisted Image Analysis I - 1TD396

Computer Exercise 1

Jyong-Jhih Lin, Linus Falk, Niklas Kostrzewa, Teng-Sung Yu

November 03, 2022

Q1

Where in the image is the pixel (1,1) located and what is the graylevel value? In command window type `I(1,1)`. Did you get the same value?

Upper left corner value: 89

```
1 imshow(I);  
2 I(1,1);  
3 Return: uint8 89
```

Q2

Explain what contrast and brightness are and include figures of the histograms of the images `napoleon.png`, `napoleon light.png`, and `napoleon dark.png`. Can you tell from the histograms which figure has the highest contrast and which figure is the brightest?

A image with high contrast will have a broader distribution of values often with peaks widely separated. While a low contrast picture will be narrower distributed. The brighter image will have its histogram further to the higher numbers in the histogram.

'napoleon.png' has the highest contrast. 'napoleon light.png' is the brightest.

```
1 listofim = ["napoleon.png", "napoleon_light.png", "napoleon_dark.png"];  
2  
3 I1 = imread(listofim(1));  
4 I2 = imread(listofim(2));  
5 I3 = imread(listofim(3));  
6  
7  
8 bins = 256;  
9 hold on;  
10  
11 subplot(3,1,1);  
12 histogram(single(I1(:)),bins)  
13 title('High contrast: ' + listofim(1))  
14 xlim([0 256])  
15  
16 subplot(3,1,2)  
17 histogram(single(I2(:)),bins)  
18 xlim([0 256])  
19 title('Brighter: ' + listofim(2))  
20  
21 subplot(3,1,3)  
22 histogram(single(I3(:)),bins)  
23 xlim([0 256])  
24 title('Darker: ' + listofim(3))
```

Q3

Explain the difference of `imagesc((I/64)*64)` and `imagesc((Is/64)*64)`, where `I` is `uint8` and `Is` is `single`.

`uint8` is an integer datatype, and its value is between 0 256. If you divided it by 64, there will be only 5 possible value (0 4) for each pixel. Therefore, after multiple by 64, there are still only five kind of values, so it will result in a vague image. `Single` is a float datatype, when it divided by 64, you retain all unique numbers, since the precision is a lot higher.




Figure 1: image of uint8 datatype




Figure 2: image of single datatype

Q4

Demonstrate a mathematical expression involving I to make it brighter.

Q4 we add a constant $g(x,y) = f(x,y) + C$, where C is a positive value.

```

1 %Q4
2 subplot(2,1,1)
3 imshow(I + 100)
4 title('I+100')
5
6 %Q5
7 subplot(2,1,2)
8 imshow(I * 0.5)
9 title('I * 0.5')
```

Q5

Demonstrate a mathematical expression involving I to give it lower contrast.

Q5 lower contrast multiply with $0 < C < 1$, $g(x,y) = f(x,y) \times C$.

1 Q6

Use $g = 2$ and $g = 12$ and explain the resulting images.

Gamma function(2) generates a overall darker image. By raise the image to the power of chosen gamma/lambda and then normalizing it to the correct range. We can see the cumulative histogram start of with a steep slope at the start indicating that the number of pixels with low value is higher then with low.

Gammafunction(0.5) generates an overall brighter image were the steep part of the image is pushed further to the higher values, right side of the cumulative histogram.




Figure 3: Example of caption

```
1 listofim = ["napoleon.png", "napoleon_light.png", "napoleon_dark.png"];
2
3 I = imread(listofim(1));
4
5 L = double(I).^2;
6 out1 = uint8(L .* (255/max(max(L))));
7 L = double(I).^0.5;
8 out2 = uint8(L .* (255/max(max(L))));
9
10 figure(3);
11 I = imread('napoleon.png');
12 subplot(3,3,1)
13 imhist(I);
14 title('original')
15 subplot(3,3,2)
16 imhist(out1);
17 title('(I).^2')
18 subplot(3,3,3)
19 imhist(out2);
20 title('(I).^0.5')
21 subplot(3,3,4)
22 imshow(I)
23 title('original')
24 subplot(3,3,5)
25 imshow(out1)
26 title('(I).^2')
27 subplot(3,3,6)
28 imshow(out2)
29 title('(I).^0.5')
30 subplot(3,3,7)
31 histogram(I,'Normalization','cdf');
32 subplot(3,3,8)
33 histogram(out1,'Normalization','cdf');
34 subplot(3,3,9)
35 histogram(out2,'Normalization','cdf');
```

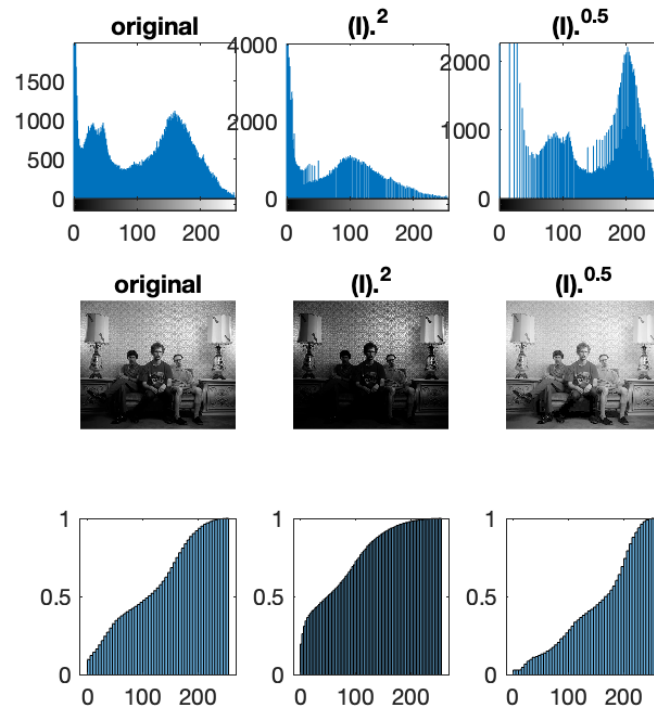


Figure 4: Example of caption

Q7

Explain how histogram equalization works in theory. Include histograms of one of the images before and after applying equalization in your report and explain what you see. Do the changes to the histograms and the images agree with the theory of histogram equalization?

Histogram equalization is a way to distribute the intensity values in a more uniform way in an image by using the cumulative density function. Perfectly evenly distributed intensities are represented by a straight line in the cumulative distribution.

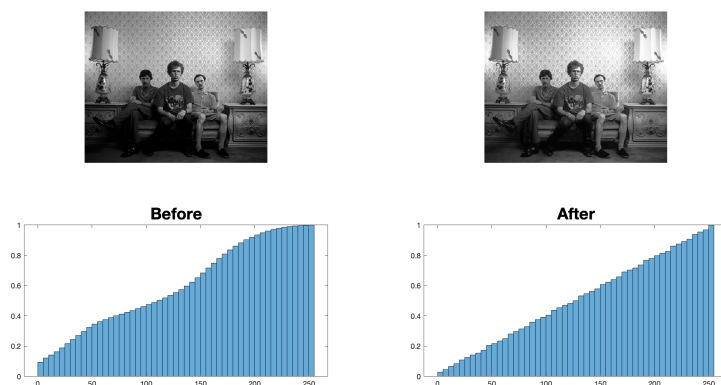


Figure 5

Q8

Explain the role of the interpolation method as well as the role of the lowpass-filter. Which combination of options do you prefer? And why? Interpolation is used when resizing the image to make it larger or fill in gaps when a image is rotated. The new pixels must be assigned values and this is done using interpolation methods: nearest or bilinear. Nearest means that we assign the new pixel the value of the nearest point it falls within, (better explained with 6) Bilinear interpolation uses the weighted average of the nearest 2×2 pixels.

Add example picture from lab?




Figure 6: https://en.wikipedia.org/wiki/Nearest-neighbor_interpolation#/media/File:Comparison_of_1D_and_2D_interpolation.svg

fix reference

Anti aliasing filter or lowpass filter is used to remove higher frequency content of the image to avoid that this information corrupts the output image after under sampling it.

Q9

Can you give a real-life example of aliasing outside the area of image analysis and signal processing in general? Have you seen this phenomenon before?

Striped sweater while having a video/conference call.

Q10

How can a “standard” healthy brain, or a mean image, of the two images brain1.png and brain2.png be constructed? In your report include a figure showing the standard brain. The mean can be created by adding the images together pixel wise and divide by the total number of images. In this case we add two images together and therefore divide by 2. One must be careful with the addition because of the data type the image is stored in. Addition easily results in “overflow”. Therefore dividing by the number of images and then add them.

```
1 Brain1 = imread('brain1.png');
2 Brain2 = imread('brain2.png');
3 Brain3 = imread('brain3.png');
4
5 standard_brain = imadd(single(Brain1), single(Brain2));
6 standard_brain = standard_brain./2;
7 figure
8 imshow(uint8(standard_brain))
```




Figure 7: standard brain upper left

2 Q11

Find the difference between the “standard” brain and the image from the stroke patient (brain3.png). Where in the brain is the change located?

Upper right corner of the image? Right frontal lobe

Q12

What happens when a pixel gets a value less than 0 or a value greater than 255? Are there other ways this can be handled? Did you think of this when you computed the “standard” brain in the previous exercise?

It becomes/stay zero if you subtract such that it would become negative, the same applies the other way: stay 255 if you add more to it. LF: Did not consider this in the first try but became aware of it when plotting the added brain images without division “burned out highlights”.

Q13

Compare rotations performed with and without interpolation. It is easiest to see differences along lines and edges of the images. What does interpolation mean in this case?

When rotating a image the repositioning of the pixels makes it necessary to take some sort of average when a pixel ends up on a boarder between pixels.



Figure 8: source <https://www.cambridgeincolour.com/tutorials/image-interpolation.htm>

Q14

In general it is faster to rotate the image by a multiple of 90 degrees than by some arbitrary degree. Explain why. For this task use Matlab functions `tic` and `toc`.

It is faster because no pixel end up on a border between pixels and no mean is calculated, no interpolation.

```

1 tic;
2 for i=1:1000
3     J = imrotate(I,20);
4 end
5 result_20deg = toc
6 tic;
7 for i=1:1000
8     J = imrotate(I,90);
9 end
10 result_90deg = toc
11
12 Output:
13 result_20deg =
14     0.2485
15
16 result_90deg =
17     0.0833

```

Q15

```

1 I = imread("my_image.jpg");
2
3 imshow(I)
4
5 J = rgb2gray(I);
6
7 subplot(1,2,1)
8 imshow(I)
9 title('Original')
10 subplot(1,2,2)
11 imshow(J)
12 title('rgb2gray')
13
14 targetSize = [683 683];
15
16 r = centerCropWindow2d(size(J),targetSize);
17 Jr = imcrop(J,r);
18 Jr = imresize(Jr, [128 128], 'bilinear', 'antialiasing', true);
19
20 figure(2)
21 imshow(Jr)
22
23
24 result = Jr;
25
26 paddda = padarray(Jr,[2 2],0);
27 result = paddda;
28
29 for i=3:128
30     for j = 3:128
31         result(i,j) = mean(paddda((i-2):(i+2),(j-2):(j+2)),'all');
32     end
33 end
34
35 result = paddda(3:130, 3:130);
36 figure(3)
37 imshow(uint8(result))
38
39 subplot(1,3,1)
40 imshow(Jr)

```



```

41 title('Original', 'fontsize',22)
42 subplot(1,3,2)
43 imshow(result)
44 title('Filtered', 'fontsize',22)
45 subplot(1,3,3)
46 imshow(Jr-result)
47 title('Subtracted', 'fontsize',22)
48 print(gcf, '-dpng', 'Q15.png')

```



Figure 9: Photo: Linus Falk




Figure 10

3 Q16

```

1 InI = Jr; %Input image
2 numofpixels=size(InI,1)*size(InI,2);
3 OutI = uint8(zeros(size(InI,1),size(InI,2))); %Output image
4 freq = zeros(256,1);
5 prob = zeros(256,1);
6
7 % Calculate frequency of pixel values and the
8 % "probability of that pixel value"
9
10
11 for i=1:size(InI,1)
12     for j=1:size(InI,2)
13         val=InI(i,j);
14         freq(val+1)=freq(val+1)+1;
15         prob(val+1)=freq(val+1)/numofpixels;

```

```

16     end
17 end
18
19 % Calculate the transformation by calculating the cdf
20 %
21 sum=0;
22 no_bins=255;
23 cum = cumsum(freq);
24 probcum = cum./numofpixels;
25 output = round(probcum.*no_bins);
26
27
28 for i=1:size(InI,1)
29     for j=1:size(InI,2)
30         OutI(i,j)=output(InI(i,j)+1);
31     end
32 end
33
34 subplot(3,1,1)
35 imshow(OutI);
36 subplot(3,1,2)
37 histogram(OutI);
38 subplot(3,1,3)
39 histogram(OutI,'Normalization','cdf');
40 title('Histogram equalization');
41 print(gcf, '-dpng', 'Q16.png')

```

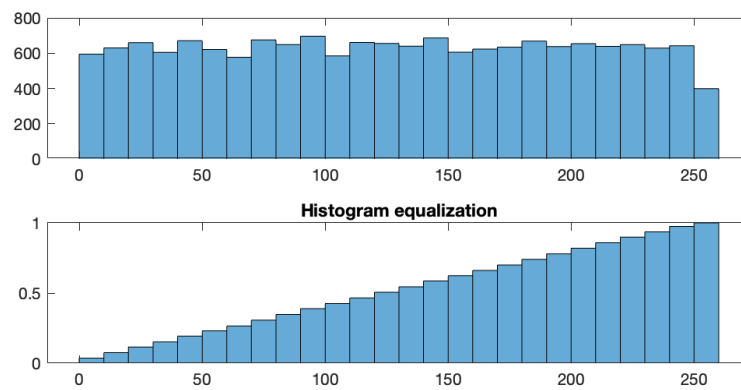


Figure 11